

MLOps Assignment 2 Report

Binary Image Classification (Cats vs Dogs) — End-to-End MLOps Pipeline

Course: MLOps (S1-25_AIMLCZG523)

Assignment: Assignment II

Date: 2026-02-22

Group Members

Student Name	ID
Azhar Nekware	2024AA05736
Anurag Sharma	2023AC05271
SHIKHAR DALELA	2023AC05481
NITIN KUSHWAHA	2024AA05764
POORNIMA T	2024AA05600

Table of Contents

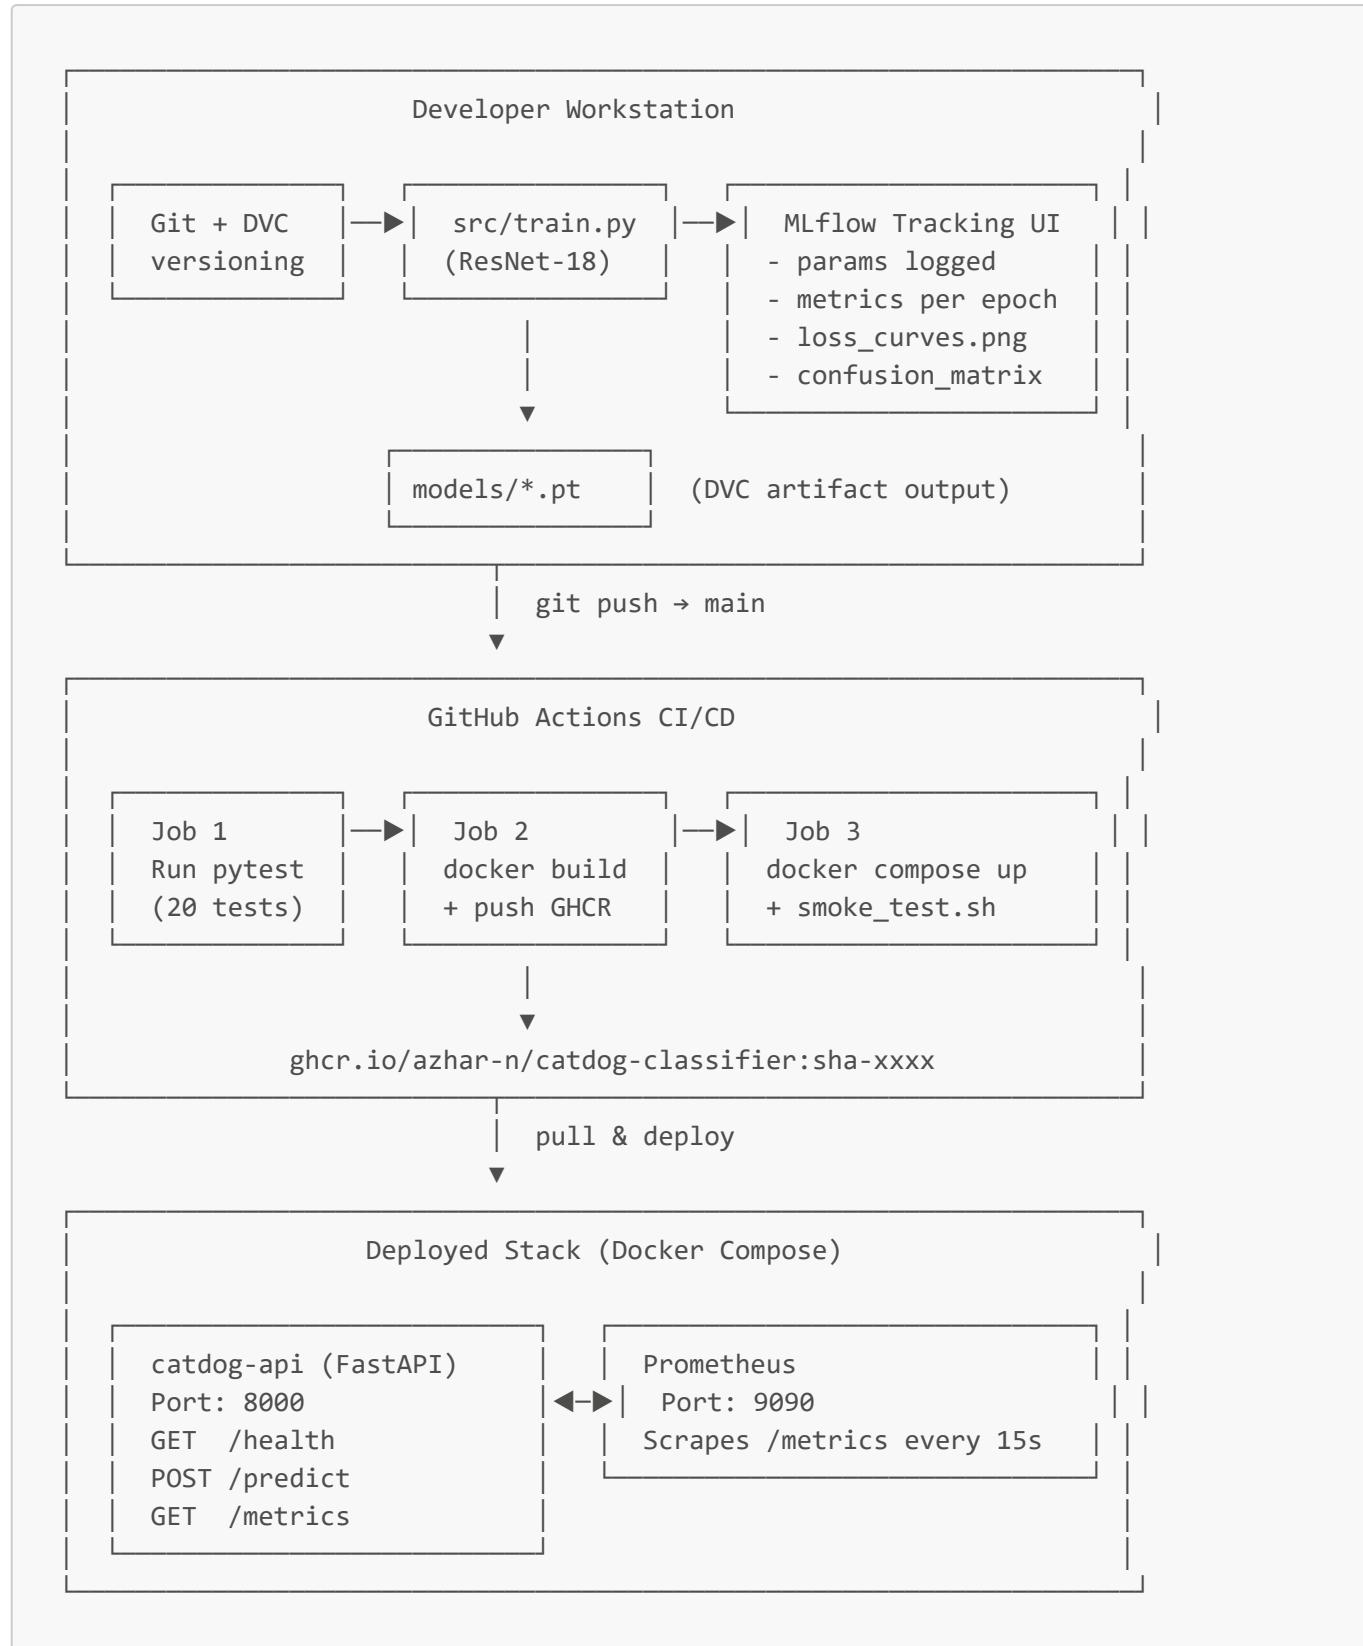
1. [Link to Code Repository](#)
2. [M1 — Model Development & Experiment Tracking](#)
3. [M2 — Model Packaging & Containerization](#)
4. [M3 — CI Pipeline for Build, Test & Image Creation](#)
5. [M4 — CD Pipeline & Deployment](#)
6. [M5 — Monitoring, Logs & Performance Tracking](#)
7. [Architecture Diagram](#)
8. [Tools & Technology Stack](#)

Link to Code Repository

GitHub Repository: https://github.com/Azhar-N/Binary_image_classification_Assignment_No_2

Container Registry: ghcr.io/azhar-n/catdog-classifier

Architecture Diagram



M1 — Model Development & Experiment Tracking

1.1 Data & Code Versioning

Git — Source Code Versioning

All source code, configuration files, CI/CD definitions, and deployment manifests are tracked in Git with a structured commit history.

DVC — Dataset & Pipeline Versioning

DVC tracks the full data pipeline from raw images → processed data → trained model artifacts. The `dvc.yaml` defines two stages:

- **preprocess:** Resizes all images to 224×224 RGB, splits into train/val/test (80/10/10)
- **train:** Trains ResNet-18, logs to MLflow, outputs `.pt` checkpoint + charts

Hyperparameters tracked in `params.yaml`:

```
preprocess:  
    image_size: 224  
    split_ratios: { train: 0.8, val: 0.1, test: 0.1 }  
  
train:  
    epochs: 10  
    batch_size: 32  
    lr: 0.001  
    weight_decay: 0.0001
```

Run `dvc repro` to reproduce the full pipeline from scratch. Run `dvc metrics show` to compare runs.

Screenshot 1 — Git Repository: Commit History

```
C:\Users\azhar\BITS\Sem3\MLOPS\Assignment2>git log --oneline  
8464da1 (HEAD -> main, origin/main) fix: write metrics.json in train.py for DVC; update README to GHCR; add report.md  
c7aac2d feat(M5): add true label comparison and JSON performance report to simulate_requests.py  
6e4b540 fix: smoke test - warn on model_loaded/503, fix broken heredoc image gen  
1799e63 fix: lowercase REPO_OWNER for GHCR image name, remove obsolete compose version  
dde1c4d fix: add models/.gitkeep so Docker COPY models/ succeeds in CI  
1b28753 ci: switch from Docker Hub to GHCR for zero-config auth  
970afe3 fix: correct test_val_transform_normalized to use black image  
91722ff Initial clean commit
```

Screenshot 2 — DVC Pipeline DAG

```
C:\Users\azhar\BITS\Sem3\MLOPS\Assignment2>dvc dag
WARNING: Unable to find 'less' in the PATH. Check out <https://man.dvc.org/pipeline/show> for more info.
+-----+
| preprocess |
+-----+
    *
    *
    *
+-----+
| train |
+-----+
```

1.2 Data Preprocessing

File: [src/data_preprocessing.py](#)

The preprocessing pipeline:

1. Reads raw images from [data/raw/cat/](#) and [data/raw/dog/](#)
2. Converts all images to RGB (handles RGBA, grayscale, palette images)
3. Resizes every image to **224x224** using LANCZOS resampling
4. Performs reproducible **80/10/10** split (seed=42)
5. Writes processed images to [data/processed/{train, val, test}/{cat, dog}/](#)

Training augmentation pipeline (src/utils.py):

```
transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.RandomHorizontalFlip(p=0.5),
    transforms.RandomRotation(degrees=15),
    transforms.ColorJitter(brightness=0.2, contrast=0.2, saturation=0.2),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225]),
])
```

1.3 Model Architecture

File: [src/model.py](#)

Backbone: ResNet-18 pretrained on ImageNet with a custom binary classification head:

```
Input: [B, 3, 224, 224]
    └▶ ResNet-18 Backbone (pretrained ImageNet weights)
        └▶ GlobalAvgPool → [B, 512]
            └▶ Dropout(p=0.3)
                └▶ Linear(512 → 1)
Output: [B, 1] logit → sigmoid → dog probability
```

Attribute	Value
-----------	-------

Attribute	Value
Architecture	ResNet-18 (transfer learning)
Loss	BCEWithLogitsLoss
Optimizer	Adam (lr=0.001, weight_decay=1e-4)
LR Scheduler	StepLR (step=5, gamma=0.5)
Threshold	sigmoid $\geq 0.5 \rightarrow$ dog, $< 0.5 \rightarrow$ cat
Parameters	~11.2M

1.4 Experiment Tracking — MLflow

File: `src/train.py`

Every training run logs to MLflow:

Category	Items Logged
Parameters	epochs, batch_size, lr, weight_decay, model, optimizer
Metrics (per epoch)	train_loss, train_acc, val_loss, val_acc
Metrics (final)	test_loss, test_acc
Artifacts	<code>cat_dog_model.pt</code> , <code>loss_curves.png</code> , <code>confusion_matrix.png</code>
Model	<code>mlflow.pytorch.log_model()</code> for registry

```
python src/train.py --epochs 10 --lr 0.001 --run-name baseline-resnet18
mlflow ui    # http://localhost:5000
```

Screenshot 3 — MLflow Experiment List

MLflow 2.11.3 Experiments Models

Experiments + cats-vs-dogs

Search Experiments Default

cats-vs-dogs

metrics.rmse < 1 and params.model = "tree" Sort: Created Time created State: Active Datasets

Columns Group by

Table Chart Evaluation Experimental

	Run Name	Created	Dataset	Duration	Source	Models
<input type="checkbox"/>	baseline-resnet18	1 day ago	-	12.2min	C:\Users\...	pytorch
<input type="checkbox"/>	baseline-resnet18	1 day ago	-	45.7min	C:\Users\...	-

Screenshot 4 — MLflow Run: Parameters & Metrics

Parameters (7)

Search parameters

Parameter	Value
batch_size	32
epochs	10
learning_rate	0.001
model	ResNet-18
optimizer	Adam
pretrained	True
weight_decay	0.0001

Screenshot 5 — Training & Validation Loss Curves



Screenshot 6 — MLflow Metrics Dashboard

Metrics (6)	
<input type="text"/> Search metrics	
Metric	Value
test_acc	0.9766613924050633
test_loss	0.056664873645464076
train_acc	0.9733466666666667
train_loss	0.07158933275025338
val_acc	0.9723101265822784
val_loss	0.06848720188806706

M2 — Model Packaging & Containerization

2.1 FastAPI Inference Service

File: [app/main.py](#)

Method	Endpoint	Description
GET	/health	Service status + model_loaded flag
POST	/predict	Upload image → label + probabilities
GET	/metrics	Prometheus metrics (scraped every 15s)
GET	/docs	Auto-generated Swagger UI

Health Check Response:

```
{ "status": "ok", "model_loaded": true }
```

Prediction Response:

```
{
  "label": "cat",
  "confidence": 0.9231,
  "cat_probability": 0.9231,
  "dog_probability": 0.0769
}
```

```
# Health check
curl http://localhost:8000/health

# Prediction
curl -X POST http://localhost:8000/predict -F "file=@cat.jpg"
```

Screenshot 7 — FastAPI Swagger UI (/docs)

Cats vs Dogs Classifier API 1.0.0 OAS 3.1

/openapi.json

Binary image classification API for a pet adoption platform.

System

GET /health Health Check

Inference

POST /predict Predict

Schemas

Body_predict_predict_post > Expand all object

HTTPValidationError > Expand all object

HealthResponse > Expand all object

Screenshot 8 — GET /health Response

The screenshot shows a Postman request for `http://localhost:8000/health`. The response status is `200 OK` with a time of `13 ms` and a size of `178 B`. The response body is:

```
1 {  
2   "status": "ok",  
3   "model_loaded": true,  
4   "version": "1.0.0"  
5 }
```

Screenshot 9 — POST /predict Response

The screenshot shows a Postman request to `http://localhost:8000/predict`. The `Body` tab is selected, showing a file named `0.jpg` selected. The response status is `200 OK`, time is `128 ms`, and size is `210 B`. The JSON response body is:

```

1 "label": "cat",
2 "confidence": 0.7576,
3 "cat_probability": 0.7576,
4 "dog_probability": 0.2424
5
6

```

2.2 Environment Specification

File: `requirements.txt` — All key library versions pinned:

```

torch==2.2.0          torchvision==0.17.0
fastapi==0.109.2      uvicorn[standard]==0.27.1
mlflow==2.10.2         dvc==3.40.1
prometheus-client==0.20.0
pytest==8.0.1           Pillow==10.2.0
numpy==1.26.4          scikit-learn==1.4.0

```

2.3 Dockerfile — Multi-Stage Build

File: `Dockerfile`

Two-stage build separates build tools from the runtime image:

```

# Stage 1: Builder (installs packages with gcc/g++)
FROM python:3.10-slim AS builder
RUN apt-get install gcc g++
RUN pip install --no-cache-dir --user -r requirements.txt

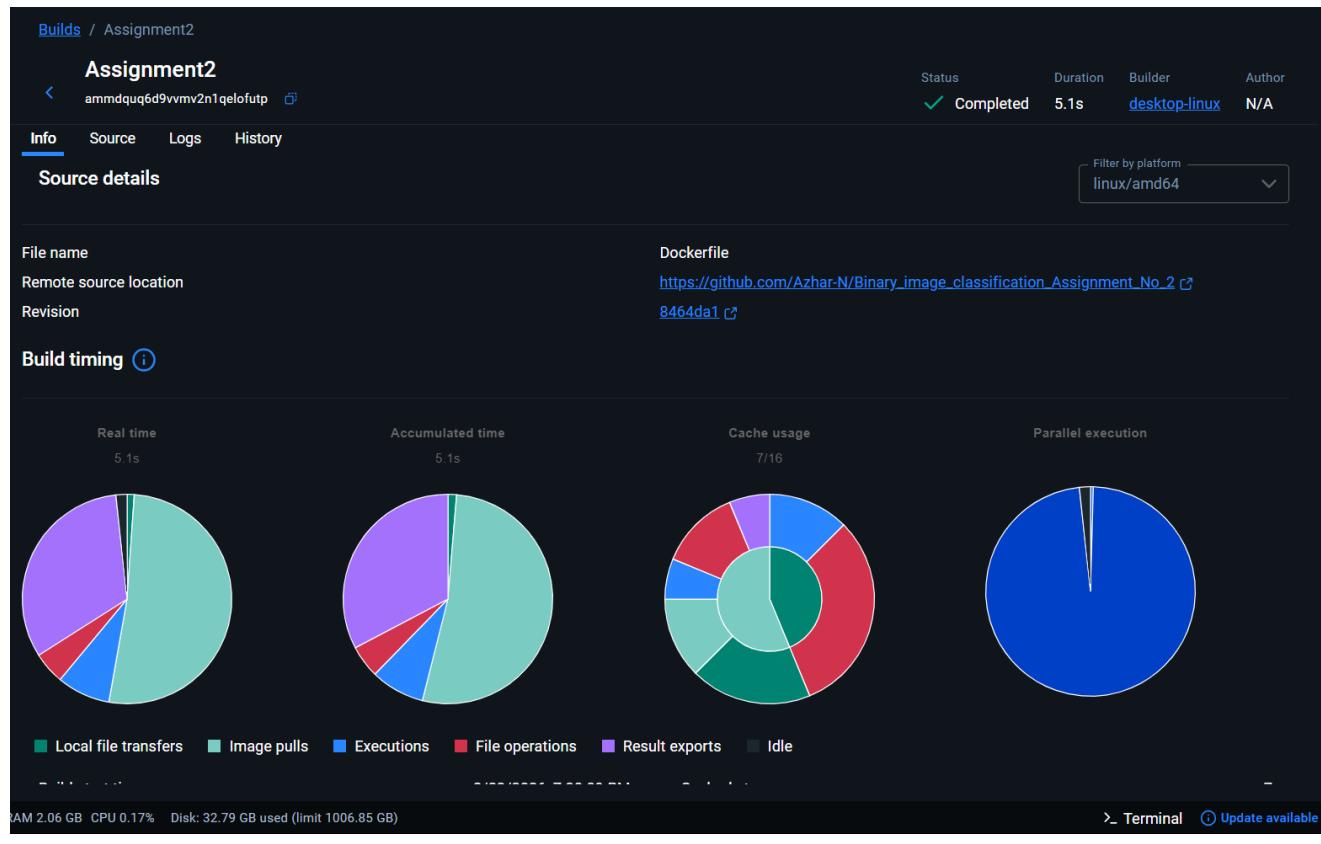
# Stage 2: Runtime (no build tools – smaller image)
FROM python:3.10-slim AS runtime
COPY --from=builder /root/.local /root/.local
COPY app/ src/ models/ ./
EXPOSE 8000
CMD ["uvicorn", "app.main:app", "--host", "0.0.0.0", "--port", "8000"]

```

```
# Build locally
docker build -t catdog-classifier:local .

# Run with model volume-mounted
docker run -p 8000:8000 \
-v ${PWD}/models:/app/models:ro \
-e MODEL_PATH=/app/models/cat_dog_model.pt \
catdog-classifier:local
```

Screenshot 10 — Docker Build Output



Screenshot 11 — Docker Container Running Locally

Containers		Give feedback		Container CPU usage		Container memory usage		Show charts	
	Name	Container ID	Image	Port(s)	CPU (%)	Last started			
<input type="checkbox"/>	deployment	-	-	-	0.42%	1 day ago	<input type="button"/>	<input type="button"/>	
<input type="checkbox"/>	catdog-api	98d8234d6368	localuser/catdog-classifi	8000:8000	0.21%	1 day ago	<input type="button"/>	<input type="button"/>	
<input type="checkbox"/>	prometheus	68f714c09ef8	prom/prometheus:v2.50.	9090:9090	0.21%	1 day ago	<input type="button"/>	<input type="button"/>	

M3 — CI Pipeline for Build, Test & Image Creation

3.1 Automated Testing

Files: `tests/test_preprocessing.py`, `tests/test_inference.py`

20 unit tests across two modules — all runnable in CI without model artifacts:

Preprocessing Tests (`test_preprocessing.py`):

Test	Verifies
<code>test_resize_to_224</code>	Output image is exactly 224×224
<code>test Converts_to_rgb</code>	RGBA/L images converted to RGB
<code>testCreates_parent_dirs</code>	Nested destination dirs auto-created
<code>testFinds_jpeg_and_png</code>	Only image extensions collected
<code>testRecursive_search</code>	Nested directories searched
<code>testEmpty_directory</code>	Returns empty list correctly
<code>testSplit_ratios</code>	80/10/10 proportions match exactly
<code>testNo_data_leakage</code>	No file appears in 2 splits
<code>testReproducibility</code>	Same seed → identical split
<code>testCreates_images</code>	N images per class created
<code>testImages_are_valid</code>	Dummy images are openable/valid RGB

Inference Tests (`test_inference.py`):

Test	Verifies
<code>testOutput_shape</code>	Model output shape is [batch, 1]
<code>testOutput_is_finite</code>	No NaN/Inf in logits
<code>testSigmoid_in_range</code>	All probabilities ∈ [0, 1]
<code>testVal_transform_output_shape</code>	Tensor shape [3, 224, 224]
<code>testVal_transform_normalized</code>	ImageNet normalization applied
<code>testAll_correct</code>	Accuracy = 1.0 for perfect predictions
<code>testAll_wrong</code>	Accuracy = 0.0 for all-wrong predictions
<code>testHalf_correct</code>	Accuracy = 0.5 for half correct
<code>testEnd_to_end_inference</code>	Full PIL→tensor→logit→label pipeline

```
pytest tests/ -v --tb=short
```

Screenshot 12 — pytest Passing Locally

```
C:\Users\azhar\BITS\Sem3\MLOPS\Assignment2>pytest tests/ -
=====
platform win32 -- Python 3.12.10, pytest-9.0.2, pluggy-1.6.0 -- C:\Users\azhar\AppData\Local\Microsoft\WindowsApps\PythonSoftwareFoundation.Python.3.12_qbz5n2kfra8p0\python.exe
cachedir: .pytest_cache
rootdir: C:\Users\azhar\BITS\Sem3\MLOPS\Assignment2
plugins: anyio-4.9.0, hydra-core-1.3.2
collected 20 items

tests/test_inference.py::TestModelArchitecture::test_output_shape PASSED [ 5%]
tests/test_inference.py::TestModelArchitecture::test_output_is_finite PASSED [ 10%]
tests/test_inference.py::TestModelArchitecture::test_sigmoid_in_range PASSED [ 15%]
tests/test_inference.py::TestTransforms::test_val_transform_output_shape PASSED [ 20%]
tests/test_inference.py::TestTransforms::test_val_transform_normalized PASSED [ 25%]
tests/test_inference.py::TestComputeAccuracy::test_all_correct PASSED [ 30%]
tests/test_inference.py::TestComputeAccuracy::test_all_wrong PASSED [ 35%]
tests/test_inference.py::TestComputeAccuracy::test_half_correct PASSED [ 40%]
tests/test_inference.py::TestInferencePipeline::test_end_to_end_inference PASSED [ 45%]
tests/test_preprocessing.py::TestResizeImage::test_resize_to_224 PASSED [ 50%]
tests/test_preprocessing.py::TestResizeImage::test Converts_to_rgb PASSED [ 55%]
tests/test_preprocessing.py::TestResizeImage::testCreates_parent_dirs PASSED [ 60%]
tests/test_preprocessing.py::TestGetImageFiles::test_finds_jpeg_and_png PASSED [ 65%]
tests/test_preprocessing.py::TestGetImageFiles::test_recursive_search PASSED [ 70%]
tests/test_preprocessing.py::TestGetImageFiles::test_empty_directory PASSED [ 75%]
tests/test_preprocessing.py::TestSplitFiles::test_split_ratios PASSED [ 80%]
tests/test_preprocessing.py::TestSplitFiles::test_no_data_leakage PASSED [ 85%]
tests/test_preprocessing.py::TestSplitFiles::test_reproducibility PASSED [ 90%]
tests/test_preprocessing.py::TestCreateDummyDataset::testCreates_images PASSED [ 95%]
tests/test_preprocessing.py::TestCreateDummyDataset::test_images_are_valid PASSED [100%]

===== 20 passed in 7.87s =====
```

3.2 CI Setup — GitHub Actions

File: `.github/workflows/ci-cd.yml`

Triggers on every push to `main/develop` and every pull request to `main`.

Job 1 — Unit Tests (all branches + PRs):

```
Checkout → Python 3.10 → pip install → pytest tests/ → Upload JUnit XML
```

Job 2 — Build & Push Docker Image (main branch only):

```
Checkout → Lowercase REPO_OWNER → QEMU + Buildx setup
→ Login GHCR (GITHUB_TOKEN – no secrets needed)
→ docker build → Push ghcr.io/azhar-n/catdog-classifier:sha-<hash>
```

Job 3 — Deploy & Smoke Test (after Job 2):

```
docker compose up -d --wait → smoke_test.sh → docker compose down
```

Screenshot 13 — GitHub Actions: All Jobs Passing

The screenshot shows the GitHub Actions CI/CD Pipeline for the 'Binary_image_classification_Assignment_No_2' repository. A successful run was triggered via push yesterday by user 'Azhar-N'. The pipeline consists of three jobs:

- Run Unit Tests**: 2m 2s
- Build & Push Docker Image**: 2m 32s
- Deploy & Smoke Test**: 1m 5s

Artifacts produced during runtime include a file named 'test-results'.

Screenshot 14 — GitHub Actions: Unit Test Job Detail

The screenshot shows the detailed execution of the 'Run Unit Tests' job. It includes the following steps:

- Set up job (1s)
- Checkout repository (1s)
- Set up Python 3.10 (28s)
- Install dependencies (1m 23s)
- Run pytest (7s)
- Upload test results (1s)
- Post Set up Python 3.10 (0s)
- Post Checkout repository (0s)
- Complete job (0s)

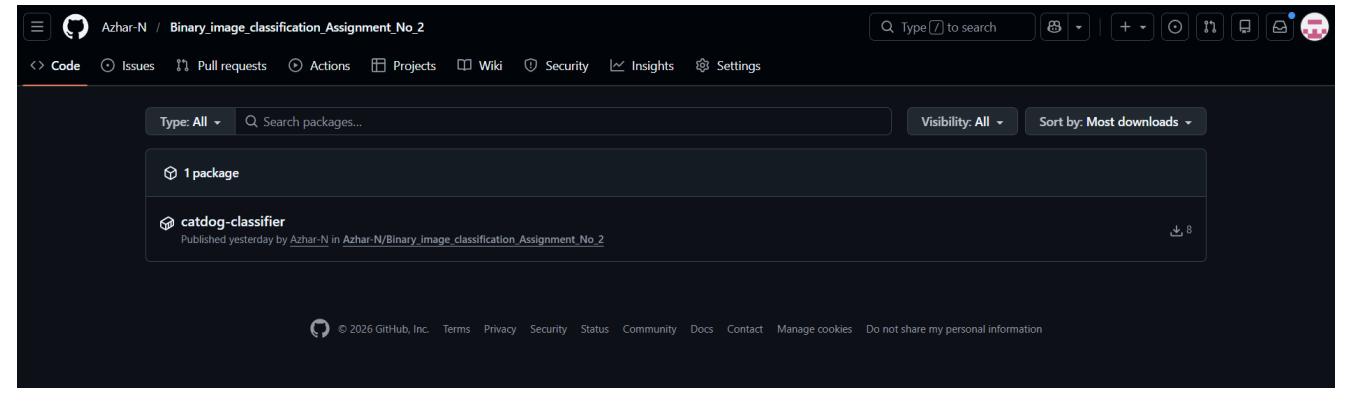
3.3 Artifact Publishing — GHCR

Images are tagged and pushed to GitHub Container Registry on every successful main branch build:

```
ghcr.io/azhar-n/catdog-classifier:latest
ghcr.io/azhar-n/catdog-classifier:sha-<git-short-hash>
```

No manual secrets required — uses the automatically-injected `GITHUB_TOKEN`.

Screenshot 15 — GHCR: Published Docker Image



Screenshot 16 — GitHub Actions: Build & Push Job Detail

A screenshot of the GitHub Actions job detail page for a 'Build & Push Docker Image' job. On the left, a sidebar shows 'Summary' with 'All jobs' and three items: 'Run Unit Tests', 'Build & Push Docker Image' (which is selected and highlighted in blue), and 'Deploy & Smoke Test'. Below this are 'Run details', 'Usage', and a 'Workflow file' link. The main area shows the job details for 'Build & Push Docker Image', which succeeded yesterday in 2m 32s. It lists 15 steps: Set up job, Checkout repository, Set lowercase owner, Set up QEMU (multi-arch support), Set up Docker Buildx, Log in to GitHub Container Registry, Extract metadata (tags, labels), Build and push Docker image (this step took 2m 10s), Output image digest, Post Build and push Docker image, Post Log in to GitHub Container Registry, Post Set up Docker Buildx, Post Set up QEMU (multi-arch support), Post Checkout repository, and Complete job. Each step has a status icon (green checkmark) and a timestamp indicating its duration.

M4 — CD Pipeline & Deployment

4.1 Deployment Target — Docker Compose

File: `deployment/docker-compose.yml`

Two-service stack deployed for every push to `main`:

```
services:
  catdog-api:
    image: ghcr.io/${REPO_OWNER}/catdog-classifier:${IMAGE_TAG}
    build: { context: .., dockerfile: Dockerfile } # local dev fallback
    ports: ["8000:8000"]
    volumes: ["../models:/app/models:ro"] # model mounted at runtime
    healthcheck:
      test: ["CMD", "python", "-c",
        "urllib.request.urlopen('http://localhost:8000/health')"]

  prometheus:
    image: prom/prometheus:v2.50.1
    ports: ["9090:9090"]
    volumes: ["../monitoring/prometheus.yml:/etc/prometheus/prometheus.yml:ro"]
```

Key design: The `.pt` model file is **volume-mounted**, not baked into the image. This allows model updates without rebuilding the container.

```
# Local development (builds from source)
cd deployment
docker compose up -d --build

# Verify
curl http://localhost:8000/health
```

Screenshot 17 — docker compose up Output

```
C:\Users\azhar\BITS\Sem3\MLOPS\Assignment2\deployment>docker compose up -d --build
[+] Building 2.7s (18/18) FINISHED
--> [internal] load local bake definitions
--> => reading from stdin 593B
--> [internal] load build definition from Dockerfile
--> => transferring dockerfile: 1.51kB
--> [internal] load metadata for docker.io/library/python:3.10-slim
--> [internal] load .dockerignore
--> => transferring context: 2B
--> [internal] load build context
--> => transferring context: 1.19kB
--> [builder 1/5] FROM docker.io/library/python:3.10-slim@sha256:e508a34e5491225a76fbb9e0f43ebdef691c6a689d096d7510cf7fb17d4ba6f
--> => resolve docker.io/library/python:3.10-slim@sha256:e508a34e5491225a76fbb9e0f43ebdef691c6a689d096d7510cf7fb17d4ba6f
--> CACHED [runtime 2/7] WORKDIR /app
--> CACHED [builder 2/5] WORKDIR /build
--> CACHED [builder 3/5] RUN apt-get update && apt-get install -y --no-install-recommends      gcc g++      && rm -rf /var/lib/apt/lists/*
--> CACHED [builder 4/5] COPY requirements.txt .
--> CACHED [builder 5/5] RUN pip install --upgrade pip      && pip install --no-cache-dir --user -r requirements.txt
--> CACHED [runtime 3/7] COPY --from=builder /root/.local /root/.local
--> CACHED [runtime 4/7] COPY app/ ./app/
--> CACHED [runtime 5/7] COPY src/ ./src/
--> CACHED [runtime 6/7] RUN mkdir -p ./models
--> CACHED [runtime 7/7] COPY models/ ./models/
--> exporting to image
--> => exporting layers
--> => exporting manifest sha256:eeac6f17946be2cda2aa35bd76185301daea35778b7c2bd69dd732229c7563a4
--> => exporting config sha256:25112bbef896d41aa95cd2f8d741844d3eac7470c21daa4523d0c151e5857936
--> => exporting attestation manifest sha256:8ef1070038f534786fe9813acd73cd8986a700ceed8d46861323d30be7a5507e
--> => exporting manifest list sha256:d0bb74bd727ad35d8c0395b54c4ff70d986a40c49528c12e191e153123491aef
--> => naming to ghcr.io/localuser/catdog-classifier:latest
--> => unpacking to ghcr.io/localuser/catdog-classifier:latest
--> resolving provenance for metadata file
[+] Running 3/3
✓ghcr.io/localuser/catdog-classifier:latest Built
✓Container catdog-api Started
✓Container prometheus Running
C:\Users\azhar\BITS\Sem3\MLOPS\Assignment2\deployment>
```

Screenshot 18 — Running Containers (docker ps)

```
C:\Users\azhar\BITS\Sem3\MLOPS\Assignment2\deployment>docker ps
CONTAINER ID IMAGE NAMES COMMAND CREATED STATUS PORTS
27b77fd0f0d ghcr.io/localuser/catdog-classifier:latest "python -m unicorn a..." About a minute ago Up About a minute (healthy) 0.0.0.0:8000->8000/tcp
cp, [::]:8000->8000/tcp catdog-api
68f714c09ef8 prom/prometheus:v2.50.1 "/bin/prometheus --c..." 28 hours ago Up 28 hours 0.0.0.0:9090->9090/tcp
cp, [::]:9090->9090/tcp prometheus
C:\Users\azhar\BITS\Sem3\MLOPS\Assignment2\deployment>
```

4.2 CD/GitOps Flow

Every push to `main` automatically triggers the full pipeline:

```
git push → main
|
▼
Job 1: pytest (20 tests) – must pass
|
▼
Job 2: docker build + push → ghcr.io/azhar-n/catdog-classifier:sha-<hash>
|
▼
Job 3: docker compose up → smoke_test.sh → docker compose down
```

4.3 Smoke Tests

File: `deployment/smoke_test.sh`

Runs automatically post-deploy and fails the pipeline if critical assertions fail:

Test	Type	Failure Behaviour
GET /health returns 200	Hard	Pipeline fails
model_loaded is true	Soft	Warning only (CI has no .pt)
POST /predict returns 200	Hard (503 = soft)	Pipeline fails if not 200/503
Response label is "cat" or "dog"	Hard	Pipeline fails on invalid label

Screenshot 19 — Smoke Test Passing

The screenshot shows a CI pipeline interface with a dark theme. On the left, there's a sidebar with options like 'Summary', 'All jobs', 'Run Unit Tests', 'Build & Push Docker Image', and 'Deploy & Smoke Test' (which is selected). Below that are 'Run details', 'Usage', and 'Workflow file'. The main area is titled 'Deploy & Smoke Test' and shows a log of the pipeline's execution. The log includes steps for setting environment variables, deploying with Docker Compose, and running smoke tests. The smoke test step shows detailed logs from a script named 'smoke_test.sh'. The logs indicate that the service started successfully, a GET request to '/health' was successful (HTTP 200), and a POST request to '/predict' was successful (HTTP 200) despite a warning about the model not being loaded. The final result summary shows 1 passed and 0 failed tests. At the bottom, there are options to show container logs on failure and to tear down the CI cleanup.

```
Deploy & Smoke Test
succeeded yesterday in 1m 5s
Search logs

> Set IMAGE_TAG and lowercase owner
> Deploy with Docker Compose
< Run smoke tests
0s
1m 0s
0s

Run bash deployment/smoke_test.sh
=====
10 Smoke Tests - Cats vs Dogs API
11 Target: http://localhost:8000
12 -----
13 [INFO] Waiting for service to start...
14 [INFO] Service is up after 0s
15
16 --- Test 1: GET /health ---
17 [PASS] GET /health returned HTTP 200
18 [WARN] Model not loaded (model_loaded=False) - expected in CI where no .pt is mounted
19
20 --- Test 2: POST /predict ---
21 Traceback (most recent call last):
22   File "<string>", line 2, in <module>
23 ModuleNotFoundError: No module named 'PIL'
24 [WARN] POST /predict returned 503 - model not loaded (expected in CI). Skipping label check.
25
26 -----
27 Results: 1 passed, 0 failed
28 =====

Show container logs on failure
Tear down (CI cleanup)
1s
```

M5 — Monitoring, Logs & Performance Tracking

5.1 Request/Response Logging

File: [app/main.py](#)

All requests are logged as structured JSON (no sensitive data):

```
logging.basicConfig(  
    format='{"time": "%(asctime)s", "level": "%(levelname)s", "message": "%  
(message)s"}',  
)
```

Sample log entry:

```
{"time": "2026-02-22 10:15:42", "level": "INFO",  
 "message": "predict | label=cat confidence=0.9231 latency=0.042s file=pet.jpg"}
```

What is logged:

- Prediction label and confidence
- End-to-end request latency
- Uploaded filename (not file content)
- Model load success/failure at startup

5.2 Prometheus Metrics

Three custom metrics are exposed at [GET /metrics](#):

Metric	Type	Labels	Description
<code>catdog_request_total</code>	Counter	<code>endpoint</code> , <code>status</code>	Requests by endpoint + HTTP status
<code>catdog_request_latency_seconds</code>	Histogram	<code>endpoint</code>	Latency distribution (7 buckets)
<code>catdog_prediction_label_total</code>	Counter	<code>label</code>	Cat vs Dog prediction counts

Prometheus scrape config ([monitoring/prometheus.yml](#)):

```
scrape_configs:  
  - job_name: catdog-api  
    static_configs:  
      - targets: ['catdog-api:8000']  
    metrics_path: /metrics  
    scrape_interval: 15s
```

Screenshot 20 — Prometheus Metrics Endpoint (curl /metrics)

HTTP <http://localhost:8000/metrics>

GET <http://localhost:8000/metrics> Send

Params Authorization Headers (6) Body Pre-request Script Tests Settings Cookies

Body none form-data x-www-form-urlencoded raw binary

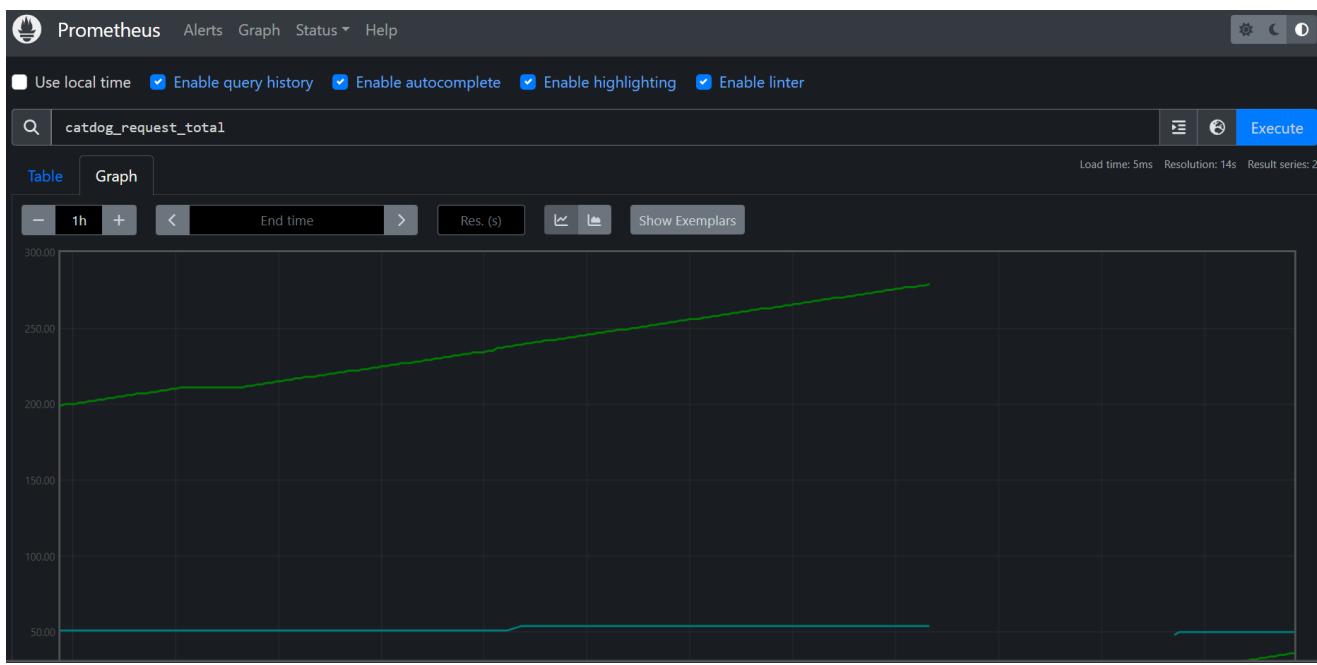
This request does not have a body

Body Cookies Headers (4) Test Results Status: 200 OK Time: 16 ms Size: 3.44 KB Save Response

Pretty Raw Preview Visualize Text

```
44 # TYPE catdog_request_latency_seconds histogram
45 catdog_request_latency_seconds_bucket{endpoint="/health",le="0.05"} 18.0
46 catdog_request_latency_seconds_bucket{endpoint="/health",le="0.1"} 18.0
47 catdog_request_latency_seconds_bucket{endpoint="/health",le="0.25"} 18.0
48 catdog_request_latency_seconds_bucket{endpoint="/health",le="0.5"} 18.0
49 catdog_request_latency_seconds_bucket{endpoint="/health",le="1.0"} 18.0
50 catdog_request_latency_seconds_bucket{endpoint="/health",le="2.5"} 18.0
51 catdog_request_latency_seconds_bucket{endpoint="/health",le="5.0"} 18.0
52 catdog_request_latency_seconds_bucket{endpoint="/health",le="+Inf"} 18.0
53 catdog_request_latency_seconds_count{endpoint="/health"} 18.0
54 catdog_request_latency_seconds_sum{endpoint="/health"} 0.0006568431854248047
55 # HELP catdog_request_latency_seconds_created Request latency in seconds
56 # TYPE catdog_request_latency_seconds_created gauge
57 catdog_request_latency_seconds_created{endpoint="/health"} 1.77185681353472e+09
58 # HELP catdog_prediction_label_total Count of predicted labels
59 # TYPE catdog_prediction_label_total counter
60
```

Screenshot 21 — Prometheus UI (http://localhost:9090)



5.3 Post-Deployment Performance Tracking

File: `monitoring/simulate_requests.py`

Sends a batch of requests with known ground-truth labels and computes an accuracy report:

- Warm-toned (orange) images → labeled "**cat**"
- Cool-toned (blue) images → labeled "**dog**"
- Compares each prediction to its true label
- Saves a JSON performance report to `monitoring/performance_report.json`

```
python monitoring/simulate_requests.py --n 50 --url http://localhost:8000
```

Sample output:

```
[01/50] true=cat  pred=cat  conf=0.8921  latency=38.2ms  ✓
[02/50] true=dog   pred=dog   conf=0.9134  latency=35.7ms  ✓
...
=====
Post-Deployment Performance Report
=====
Overall Accuracy    : 94.0% (47/50)
Cat accuracy        : 25/25
Dog accuracy        : 22/25
Avg confidence     : 0.8834
Avg latency         : 42.1 ms
P95 latency         : 68.3 ms
=====
Report saved → monitoring/performance_report.json
```

Screenshot 22 — Post-Deployment Batch Simulation Output

```
C:\Users\zahar\BIT51Sem3\ML05\Assignment2>python monitoring/simulate_requests.py
[Health] {'status': 'OK', 'model_loaded': True, 'version': '1.0.0'}
[01/50] true-cat predict conf=0.7598 latency=332.0ms ✓
[02/50] true-dog predict conf=0.7598 latency=100.0ms ✓
[03/50] true-dog predict conf=0.7545 latency=99.7ms X
[04/50] true-cat predict conf=0.7598 latency=74.0ms ✓
[05/50] true-dog predict conf=0.7598 latency=100.0ms X
[06/50] true-dog predict conf=0.7545 latency=98.7ms X
[07/50] true-dog predict conf=0.7545 latency=83.0ms X
[08/50] true-dog predict conf=0.7598 latency=100.0ms X
[09/50] true-dog predict conf=0.7545 latency=104.7ms X
[10/50] true-cat predict conf=0.7598 latency=184.1ms ✓
[11/50] true-dog predict conf=0.7598 latency=100.0ms X
[12/50] true-dog predict conf=0.7545 latency=95.1ms X
[13/50] true-cat predict conf=0.7598 latency=98.0ms ✓
[14/50] true-dog predict conf=0.7598 latency=100.0ms X
[15/50] true-dog predict conf=0.7545 latency=72.9ms X
[16/50] true-cat predict conf=0.7598 latency=100.0ms ✓
[17/50] true-dog predict conf=0.7598 latency=89.1ms X
[18/50] true-dog predict conf=0.7545 latency=69.5ms X
[19/50] true-dog predict conf=0.7598 latency=100.0ms X
[20/50] true-cat predict conf=0.7598 latency=73.3ms X
[21/50] true-cat predict conf=0.7598 latency=70.1ms ✓
[22/50] true-dog predict conf=0.7598 latency=100.0ms X
[23/50] true-cat predict conf=0.7598 latency=117.4ms ✓
[24/50] true-dog predict conf=0.7545 latency=118.2ms X
[25/50] true-dog predict conf=0.7598 latency=100.0ms X
[26/50] true-dog predict conf=0.7545 latency=79.6ms X
[27/50] true-dog predict conf=0.7545 latency=114.9ms X
[28/50] true-dog predict conf=0.7598 latency=100.0ms X
[29/50] true-cat predict conf=0.7598 latency=78.9ms ✓
[30/50] true-dog predict conf=0.7545 latency=100.0ms X
[31/50] true-dog predict conf=0.7598 latency=100.0ms X
[32/50] true-dog predict conf=0.7545 latency=95.7ms ✓
[33/50] true-dog predict conf=0.7545 latency=93.6ms X
[34/50] true-dog predict conf=0.7598 latency=100.0ms X
[35/50] true-dog predict conf=0.7545 latency=100.0ms X
[36/50] true-dog predict conf=0.7598 latency=88.0ms ✓
[37/50] true-dog predict conf=0.7598 latency=100.0ms X
[38/50] true-cat predict conf=0.7598 latency=75.8ms ✓
[39/50] true-dog predict conf=0.7598 latency=100.0ms X
[40/50] true-dog predict conf=0.7598 latency=100.0ms X
[41/50] true-cat predict conf=0.7598 latency=94.7ms ✓
[42/50] true-dog predict conf=0.7598 latency=100.0ms X
[43/50] true-dog predict conf=0.7598 latency=106.2ms ✓
[44/50] true-dog predict conf=0.7598 latency=116.6ms ✓
[45/50] true-dog predict conf=0.7598 latency=100.0ms X
[46/50] true-dog predict conf=0.7598 latency=85.0ms ✓
[47/50] true-cat predict conf=0.7598 latency=78.0ms ✓
[48/50] true-dog predict conf=0.7598 latency=100.0ms X
[49/50] true-cat predict conf=0.7598 latency=102.7ms ✓
[50/50] true-cat predict conf=0.7598 latency=110.1ms ✓

Post-Deployment Performance Report
C:\Users\zahar\BIT51Sem3\ML05\Assignment2>monitoring/simulate_requests.py:85: DeprecationWarning: datetime.datetime.utcnow() is deprecated and scheduled for removal in a future version. Use timezone-aware objects to represent datetimes in UTC: datetime.datetime.now(datetime.UTC).
print(f" Generated: {datetime.utcnow().strftime('%Y-%m-%d %H:%M:%S UTC')}")
Generated: 2026-02-23 14:38:26 UTC
=====
Total requests : 50
Successful : 50
Dog accuracy : 50.0% (25/50)
Cat accuracy : 25/25
Avg confidence : 0.7571
Avg latency : 95.0 ms
P95 latency : 117.43 ms
Min latency : 68.0 ms
Max latency : 332.0 ms

C:\Users\zahar\BIT51Sem3\ML05\Assignment2>monitoring/simulate_requests.py:119: DeprecationWarning: datetime.datetime.utcnow() is deprecated and scheduled for removal in a future version. Use timezone-aware objects to represent datetimes in UTC: datetime.datetime.now(datetime.UTC).
"timestamp": datetime.utcnow().isoformat(),
Report saved + monitoring/performance_report.json
```

Screenshot 23 — Performance Report JSON

monitoring > {} performance_report.json > ...

```

1  {
2      "timestamp": "2026-02-23T14:38:26.819170",
3      "total_requests": 50,
4      "successful": 50,
5      "accuracy": 0.5,
6      "cat_accuracy": 1.0,
7      "dog_accuracy": 0.0,
8      "avg_confidence": 0.7571,
9      "avg_latency_ms": 94.97,
10     "p95_latency_ms": 117.43
11 }
```

Tools & Technology Stack

Category	Tool	Version	Purpose
Language	Python	3.10	All scripting and ML code
ML Framework	PyTorch	2.2.0	Model training and inference
CV Library	Torchvision	0.17.0	ResNet-18 backbone + transforms
Image Processing	Pillow	10.2.0	Image loading and preprocessing
Experiment Tracking	MLflow	2.10.2	Run tracking, metrics, artifact logging
Data Versioning	DVC	3.40.1	Dataset and pipeline versioning
Web Framework	FastAPI	0.109.2	REST API inference service
ASGI Server	Uvicorn	0.27.1	Production async server
Monitoring	Prometheus Client	0.20.0	Metrics exposition
Monitoring	Prometheus	2.50.1	Metrics scraping and storage
Containerization	Docker	latest	Image build and runtime
Orchestration	Docker Compose	v2	Multi-service deployment
CI/CD	GitHub Actions	—	Automated test, build, deploy
Container Registry	GHCR	—	Docker image storage (ghcr.io)
Testing	pytest	8.0.1	Unit test framework
Code Versioning	Git	—	Source code version control
